

The Objective Ledger™

TOL-OS · Baseline v1.0.4

Spread Love & Acceptance Relationship Survival Kit

Full Program bundle — Preface, Chapters I through V, and Workpapers 01 through 04, 09, and 13.

spreadloveandacceptance.com

Preface: Unbilled Debt

Every household runs a set of books nobody agreed to keep.

Someone notices the milk is low. Someone remembers the dentist appointment three weeks out. Someone feels the temperature of the room change before an argument starts and quietly redirects it. None of this shows up on a calendar. None of it gets a thank-you, because thank-yous require the work to be visible first, and this work is built to be invisible — that's what makes it work.

Over months and years, that invisible work accumulates like a debt that only one party can see on their books. The other party isn't necessarily refusing to pay it. They usually don't know it's owed. That gap — between labor performed and labor acknowledged — is what this book calls **unbilled debt**.

This is not a book about keeping score to win an argument. A ledger that only one person gets to read isn't a ledger, it's a grudge with better formatting. The point of writing anything down is the opposite of scorekeeping: it's making the invisible column visible to *both* people, in the same units, at the same time, so the conversation can start from shared facts instead of competing impressions.

You'll notice this book borrows vocabulary from accounting, systems engineering, and nervous-system science. That's deliberate, not decorative. These fields are good at one specific thing: describing a system's actual state without moralizing about who's to blame for it. A ledger that's out of balance isn't evil. A nervous system in a stress response isn't broken. Both are just information — a starting point, not a verdict on either person's worth.

What follows is structured as a working ledger, not a lecture. Read the chapters, use the tools, and close the books together at the end of each week. The goal isn't a perfect 50/50 split measured to the decimal. It's a household where both people can see the whole picture, and neither one is quietly carrying debt the other doesn't know exists.

Chapter I — The Radio Frequency Paradigm

A model, not a measurement. Nothing in this chapter involves recording, analyzing, or scoring anyone's actual voice — it's a mental picture borrowed from physics to explain why calm people say cruel-sounding things to each other under stress.

The heterodyne squeal

Tune two radios to frequencies that are close but not identical, and you don't get silence — you get a squeal. The two signals overlap just enough to interfere, producing a third, uglier sound that neither radio was actually broadcasting. Engineers call this a heterodyne: the audible beat frequency created by two signals that are almost, but not quite, in sync.

This book borrows that image for a familiar moment: two people who are each saying something reasonable, in a tone shaped by their own fatigue, urgency, or nervous-system state, and getting a fight out of it that neither of them intended to start. The words weren't the problem. The mismatch in frequency was.

What sets the frequency

Three things tend to shift someone's "transmission frequency" without their noticing:

- **Pace.** Someone who's been mobilized (see Chapter II's ladder — sympathetic activation) talks faster, interrupts sooner, and reads neutral pauses as stonewalling.
- **Register.** A comment meant as a status update ("the sink's full again") can be transmitted — or received — as an accusation, depending on the tone underneath it.
- **Urgency.** One person's "when you get a chance" and another person's "when you get a chance" can mean completely different windows of time.

None of this requires either person to be wrong. Two correctly-functioning radios on slightly different frequencies will squeal every time, and no amount of arguing about who's "really" at fault fixes the interference. What fixes it is retuning.

Retuning, practically

Before continuing a conversation that's started to squeal:

1. Name the mismatch out loud, without blame: *"I think we're on different frequencies right now."*
2. Check your own state against the ladder in Chapter II. You can't accurately retune while mobilized or shut down.
3. Slow the pace deliberately — shorter turns, more pauses, lower volume — until the interference clears.
4. Only then return to the actual content of the conversation.

The Radio Frequency Paradigm isn't a diagnosis of who's transmitting badly. It's a shared vocabulary for the moment when both of you are just out of sync, so neither of you has to carry the whole blame for a fight that was mostly a timing problem.

Chapter II — The Epistemic Verdict Engine & P(Solvency)

Important framing before anything else: this chapter describes a heuristic self-reflection tool, not a scientific instrument. "P(Solvency)" is not a validated statistical model and it does not sit in judgment of either person or of the relationship's worth. It scores one narrow thing — whether the current household workload arrangement looks sustainable on paper — the same way a small-business cash-flow snapshot tells you about this month's operations, not about whether the business (or its owners) are good people.

What the Verdict Engine is for

Every long-running household eventually gets asked, in one form or another: "is this actually working?" Usually the question gets answered with whichever incident happened most recently — last night's dishes, this morning's tone. That answers the wrong scope.

The Verdict Engine asks a narrower, calmer question: **based on what's actually written in the ledger this month — hours logged in the Lemonade Stand, tasks with a named owner in the RACI table, and self-reported autonomic load — does the current workload arrangement look sustainable, or is it quietly running a deficit?**

The inputs

P(Solvency) is a weighted read of three ledger inputs, each already covered elsewhere in this book:

Input	Source	What it captures
Workload balance	Lemonade Stand (Ch. III)	How evenly visible hours are distributed
Ownership clarity	RACI Treaty (Workpaper 03)	How many recurring tasks have a clear, undisputed owner
Autonomic load	Nervous System Ladder (Ch. II sidebar) / ASI (Workpaper 02)	How much of the friction is workload versus baseline stress

A plain-language formula

Rather than a black-box score, treat P(Solvency) as a simple weighted read:

```
solvency_read = (workload_balance_score * 0.4)
+ (ownership_clarity_score * 0.35)
+ (1 - average_autonomic_saturation) * 0.25
```

Each component is scored 0–1 from the workpapers, not from a mood in the moment. The output is a number between 0 and 1, and it's meant to be read as a **snapshot**, not a sentence:

- **Above ~0.7:** The current arrangement is carrying its own weight. Keep the same rhythm of check-ins.
- **0.4–0.7:** Something in the three inputs is drifting — usually ownership clarity first. Revisit the RACI table before anything else.
- **Below ~0.4:** The workload as currently structured isn't sustainable. That's a statement about the *arrangement*, not about either person's love, effort, or worth — arrangements get renegotiated all the time.

What this number is not

It is not a relationship-compatibility score. It is not a fairness verdict about who's the better partner. It does not use anyone's tone of voice, private thoughts, or communication style as an input — only the three ledger sources above, which both people fill in together. If it's ever used to end a conversation ("well, the number says...") instead of start one, it's being misused. The number's only job is to give two people a shared, low-drama starting point for the conversation Chapter V calls closing the books.

Chapter III — Autonomic Saturation & the 7 Ocular Vectors

Both tools in this chapter are self-report frameworks, filled in by the person they describe. Neither one measures anything automatically, and neither should be filled in by one partner about the other — that turns a reflection tool into a surveillance tool, which is exactly what this book is trying to avoid.

Autonomic Saturation

"Saturation" here just means: how much of your current stress response is left over from somewhere else, versus caused by the thing in front of you right now. A short version lives on the site as the Nervous System Ladder (ventral / sympathetic / dorsal). The full **Autonomic Saturation Index**, detailed in Workpaper 02, breaks that single reading into a handful of self-rated inputs — sleep, workload elsewhere, recent conflict, physical state — so that "I'm upset about the dishes" and "I'm exhausted and the dishes are what I noticed" stop getting treated as the same problem.

A high saturation reading isn't an excuse or a diagnosis. It's a note to both people: *whatever gets decided in the next ten minutes should probably get revisited later, when saturation is lower.*

The 7 Ocular Vectors

"Ocular" is used loosely here to mean *vantage point* — the angle you're looking at a friction point from. Most arguments happen because two people are each looking at the same event through a different one of these seven lenses and assuming the other person sees what they see. The seven vectors:

1. **Physical fatigue** — how rested either of you actually is right now
2. **Time scarcity** — whose calendar is genuinely fuller this week
3. **Historical pattern** — whether this specific friction has come up before
4. **Family script** — what "normal" looked like in the household either of you grew up in
5. **Fairness perception** — whether the *split* feels fair, independent of the raw hours
6. **Communication style** — direct versus indirect, high-context versus low-context
7. **External stressors** — money, work, health, anything outside the relationship itself

Before a disagreement about chores turns into a disagreement about character, it's worth naming which of these seven either of you is actually looking through. Two people can both be completely right about what they're seeing and still be looking at different vectors of the same event.

Using both together

Autonomic Saturation asks *how much charge is in the room*. The Ocular Vectors ask *what each person thinks the room is actually about*. Used before a conversation rather than during one, they replace two people guessing at each other's internal states with two people just saying what those states are.

Chapter IV — Deontological Parity & Sensory Gating

Deontological parity

Most chore arguments are actually fairness-philosophy arguments wearing a dish towel. There are two common ways to judge whether a division of labor is "fair":

- **Outcome-based (consequentialist):** did the right things get done, by whoever had time or energy, however it worked out this week?
- **Duty-based (deontological):** did each person do the specific things they agreed to own, regardless of how the week went?

Both are reasonable. The friction shows up when one person is judging the week by outcomes ("it all got done, who cares who did it") and the other is judging it by duties ("you said you'd handle trash, and you didn't, again"). Neither framework is wrong — but a household needs to agree on which one it's using, the same way a contract needs governing law.

Deontological parity is this book's term for a simple agreement: once a task has a named owner in the RACI table (Workpaper 03), it's evaluated on whether the owner did their part — not on whether the outcome happened to work out anyway because someone else quietly covered it. This protects the person who keeps covering gaps from having that coverage used as evidence that the gap doesn't matter.

Sensory gating

Borrowed loosely from neuroscience, where it describes the brain's ability to filter out repetitive or low-priority stimuli so attention isn't consumed by everything at once. Used here as a practical rule rather than a clinical claim: **not every comment needs to be let in for full processing the moment it's said.**

When you're mobilized or shutting down (Chapter II's ladder), a neutral comment can arrive at full volume and demand an immediate reaction. Sensory gating, practiced deliberately, looks like:

1. Noticing the comment landed harder than its content probably warranted.
2. Naming that out loud — *"I heard that as bigger than it probably was, give me a second"* — instead of reacting from the spike.
3. Deciding, once your state has settled, whether it actually needs a response now, later, or not at all.

This isn't suppression — the goal isn't to gate everything out permanently. It's timing: letting the nervous system finish its spike before deciding how much weight a comment actually deserves. Combined with deontological parity, it gives both people a way to separate "this task wasn't done" (a fact, evaluated against an agreed duty) from "this landed as an attack" (a reaction, evaluated against

your current state) — instead of collapsing both into the same fight.

Chapter V — The Deficit Audit

Like the tools in earlier chapters, the audit described here evaluates a shared arrangement over time — not either person's character, and not a settlement of past hours owed.

Why a monthly pass, on top of the weekly ones

By this point in the book, two rhythms are already running: the daily 90-second loop (Workpaper 13) and the weekly closing, where the Lemonade Stand gets tallied and the RACI treaty gets a quick glance. Both are useful precisely because they're small and frequent — but frequent, small checks have a blind spot. A task that goes unclaimed for one week looks like a fluke. The same task going unclaimed for four weeks in a row is a pattern, and patterns are invisible at the weekly scale; they only show up once you step back and look at a month at a time.

The Deficit Audit is that step back. Once a month, instead of asking "how did this week go," the question becomes: **which gaps kept reopening, no matter how many times we closed them?**

What a deficit actually looks like at this scale

A single missed task is an event. A deficit is a task that keeps generating the same event, month after month, without ever landing on a stable owner. It's the difference between "the trash didn't go out Tuesday" and "the trash has quietly become no one's job for three months running." The first is noise. The second is a structural gap in the RACI treaty that weekly check-ins keep patching over instead of fixing.

This matters because unresolved structural gaps are exactly where the Anchor Tax described elsewhere in this framework compounds fastest — not from any single missed task, but from the same task being renegotiated from scratch every time it resurfaces, because it never got a name attached to it that stuck.

What the audit is for, and what it isn't

The Deficit Audit, detailed as a working template in Workpaper 04, exists to do three things: surface which tasks are genuinely recurring gaps rather than one-off misses, decide whether each gap needs a new RACI owner or a capacity conversation, and then close the month cleanly so next month starts from a fixed structure instead of the same unresolved argument.

It is explicitly not a monthly performance review of either person, and it does not retroactively tally what was "owed" in hours already passed. Those hours are gone; billing for them backward doesn't fix anything and just turns an audit into a grievance. The only question the Deficit Audit asks about the past month is where the *structure* failed to hold — so that structure, not effort or affection, is what gets adjusted going into the next one.

WP-01 — Field Audit & Neutral Refusals

Purpose: capture what actually happened this week before either person's memory of it hardens into a story. Fill this in together, or independently and compare — don't fill it in *about* the other person.

Part A — Field Audit

A field audit is just a raw log, kept for 5–7 days, before any interpretation gets added. One row per task noticed or done.

Day	Task noticed / done	Who did it	Minutes (rough)	Was it asked for, or just noticed and handled?
Mon				
Tue				
Wed				
Thu				
Fri				
Sat				
Sun				

At the end of the week, total the minutes per person. This total feeds the Lemonade Stand (</index.html>, Chapter III) and the `workload_balance_score` used in the solvency calculator.

Audit rule: log what happened, not what should have happened. "I noticed the trash needed to go out three times before doing it" is one row, logged honestly, not three rows of resentment.

Part B — Neutral Refusals

Half of unbilled debt isn't unequal effort — it's one person being unable to say no to a request without it becoming a conflict. A neutral refusal is a scripted way to decline or defer a task that states the boundary without relitigating whether the request was reasonable.

Use one of these, or write your own in the same shape: **acknowledge** → **state your actual capacity** → **offer a specific alternative**.

- *"I hear that needs doing. I don't have the bandwidth for it today — can it wait until tomorrow, or is there someone else who could take it?"*

- *"I can do that, but something else has to move if I do. Which one matters more this week?"*
- *"I'm at capacity right now. I can commit to it by [specific time] if that works."*

What a neutral refusal is not: a justification essay, a counter-list of everything you already did, or a refusal to ever revisit the request. It's a single, calm boundary statement that keeps the door open for renegotiation instead of ending the conversation.

Sign-off

	Name	Date
Prepared by		
Reviewed by		

WP-02 — Autonomic Saturation Index (ASI)

Purpose: a self-report matrix, filled in by each person about themselves, that separates "how much stress am I already carrying" from "how upset am I about this specific thing." Not a clinical instrument — a structured gut-check.

Fill this in independently. Score each row 0 (not at all) to 4 (maximally true) for how the last 24–48 hours have actually gone.

Factor	0	1	2	3	4
Sleep debt (less rested than usual)					
Workload elsewhere (job, school, caregiving)					
Recent unresolved conflict (any source)					
Physical state (hungry, sick, in pain)					
Time pressure today specifically					

Average Saturation Score = sum of the five scores ÷ 20

This produces a number between 0 and 1. That number is the `average_autonomic_saturation` input used in the Verdict Engine (Chapter II) and the solvency calculator.

Reading your own number

- **0.0–0.3:** Low saturation. Whatever's coming up is probably actually about the thing itself.

- **0.3–0.6:** Moderate. Worth naming out loud before a hard conversation: *"heads up, I'm carrying more than usual today."*
- **0.6–1.0:** High. Per the Nervous System Ladder (Chapter II), this is a signal to postpone anything that doesn't need deciding in the next hour.

What this worksheet is not for

Don't fill this in about your partner, and don't use your own high score as a way to end a conversation you don't want to have — it's a postponement tool, not an exit. A high score means *"let's revisit this tomorrow,"* not *"this doesn't need to happen."*

Sign-off

	Score	Date
Partner A		
Partner B		

WP-03 — Domestic RACI Treaty

Purpose: a living agreement assigning exactly one Responsible and one Accountable name to every recurring household task, so ownership stops getting renegotiated by default every week.

- **Responsible (R):** does the task.
- **Accountable (A):** notices if it didn't happen and follows up. Can be the same person as R.

The Treaty

Task	Frequency	Responsible	Accountable	Notes
Groceries	Weekly			
Cooking	Daily			
Dishes	Daily			
Laundry	Weekly			
Bills & scheduling	Monthly			
Cleaning (bathroom/kitchen)	Weekly			
Pet care	Daily			
Car maintenance	As needed			
Social/family calendar	Ongoing			
Emotional check-ins	Weekly			
<i>(add rows as needed)</i>				

Amendment process

This treaty isn't permanent — it's meant to be renegotiated when life changes (new job, health issue, new baby, etc.), not silently reassigned when one person just starts doing more. To amend:

1. Either person can call for a review at the weekly closing (Chapter V). 2. Changes are written into the table above, not just agreed verbally. 3. Both people initial the change.

Ratification

By initialing below, both people confirm they've read the current version of this table and agree to the ownership as written — not that every task feels perfectly fair, only that ownership is clear.

	Initials	Date
Partner A		
Partner B		

Ownership Clarity Score (for the Verdict Engine, Chapter II): count tasks with both an R and an A filled in, divided by total tasks listed. This is the `ownership_clarity_score` input.

WP-04 — Unbilled Deficit Audit

Purpose: a monthly reconciliation, done together, that looks across four weeks of Field Audit logs (WP-01) and the RACI Treaty (WP-03) to find tasks that keep becoming deficits — not to tally what either person "owes," but to find where the household's structure needs a fix rather than another one-off patch.

Cadence: monthly. Not weekly — this audit is specifically for patterns that are invisible at the weekly scale.

Part A — Pull the month's raw data

From your four weekly Field Audits (WP-01), list every task that was logged as noticed-but-undone, or done only after being raised more than once, in **any** week this month.

Task	Week 1	Week 2	Week 3	Week 4	Times flagged this month
------	--------	--------	--------	--------	--------------------------

Mark a week with a check if the task was noticed as a gap that week (undone, or raised more than once before being done). A task flagged 3–4 times this month is a genuine deficit, not a fluke.

Part B — Classify each deficit

For every task flagged twice or more, decide which of these it actually is:

Task	Has a named RACI owner? (WP-03)	Classification	Action
	Y / N	Structural gap | Capacity issue | One-off, no action	
	Y / N	Structural gap | Capacity issue | One-off, no action	

- **Structural gap:** no clear owner exists, or the named owner isn't the person actually handling it in practice. Fix: update the RACI Treaty (WP-03) this session.

- **Capacity issue:** has a named owner, but that person has been consistently unable to keep up with it. Fix: a direct conversation using a neutral refusal or reassignment (WP-01, WP-09) — not a private resolution to just try harder.

- **One-off, no action:** flagged more than once but tied to an identifiable temporary cause (illness, travel, a busy work month). Note it and move on; don't let a temporary cause justify a permanent

structural change.

Part C — Close the audit

1. Update the RACI Treaty (WP-03) for every task classified as a structural gap. 2. Schedule the capacity conversation for anything classified as a capacity issue — this audit surfaces it, it doesn't replace the actual conversation. 3. Record the count of structural gaps found this month below. A falling count month over month is the clearest sign this framework is working; a rising count means something bigger changed and may need its own conversation.

Month	Structural gaps found	Capacity issues found	One-offs
-------	-----------------------	-----------------------	----------

What this audit is not for

- **Not a retroactive bill.** Hours already passed don't get tallied up and "charged" to anyone. The audit only looks backward to find the pattern, and only acts forward on the structure.
- **Not a performance review.** A task with no owner is a gap in the treaty, not a verdict on the person who happened to keep covering it.
- **Not a substitute for the weekly closing.** This audit works because the weekly rhythm (Lemonade Stand, RACI glance) keeps feeding it real data. Skipping the weekly checks makes this audit far less accurate.

Sign-off

	Initials	Date
Partner A		
Partner B		

Feeds into: the `ownership_clarity_score` used by the Verdict Engine (Chapter II) — a month with fewer unresolved structural gaps should show up as a higher score the following month.

WP-09 — Tone Transducer & Filter

Purpose: a practical checklist for "retuning" before responding, operationalizing the Radio Frequency Paradigm from Chapter I. This does not involve recording or analyzing anyone's actual voice — it's a self-check applied by each person to their own next message.

The transducer step — before you speak or send

A transducer converts one form of signal into another. Here, it means converting a raw internal reaction into something the other person can actually receive. Before responding to something that landed hard, run it through three questions:

1. **What's the fact underneath this?** (One sentence, no adjectives.)
2. **What's the feeling underneath this?** (Name it — frustrated, tired, unseen, rushed.)
3. **What's the actual ask?** (What do you want to happen next, specifically?)

Write all three down before responding, even mentally. A message built from fact → feeling → ask transmits far more cleanly than one built from the raw reaction.

The filter step — before you react to what they said

The filter runs the other direction: deciding how much weight to give an incoming message before responding to it.

Check	Yes / No
Is this about a specific, nameable task or event?	
Would I read this the same way if I weren't already at high saturation (WP-02)?	
Is there a neutral interpretation that also fits what was said?	
Am I responding to their words, or to a pattern from a past conversation?	

If saturation is high or the neutral interpretation hasn't been ruled out, use a neutral refusal or a pause script (WP-01) before continuing.

Worked example

- Raw reaction: *"You never think about anyone but yourself."*
- Fact: "The trash wasn't taken out before pickup."

- Feeling: "Frustrated, and a little unseen since I mentioned it twice."
- Ask: "Can we agree on a fixed day/time for trash so it's not something I have to keep bringing up?"

Same underlying issue, transmitted on a frequency the other person can actually tune into.

Sign-off

	Reviewed	Date
Partner A		
Partner B		

WP-13 — Phase-Locked Loop (PLL) Protocol

Purpose: a recurring, low-effort check-in ritual, named after the electronics circuit that keeps two signals in sync through constant small corrections rather than occasional large ones. Applied here as a daily habit that prevents the "frequency drift" described in Chapter I from building up into a squeal.

Why small and frequent beats big and rare

A phase-locked loop doesn't wait for two signals to drift far apart and then yank them back into sync — it makes tiny corrections continuously, so drift never accumulates. The domestic equivalent: a 90-second daily check-in prevents most of what would otherwise become a 45-minute weekend argument.

The daily loop (about 90 seconds)

Pick a consistent moment — after dinner, before bed, whatever's realistic. Each person answers, out loud, in one sentence each:

1. **Load:** "Today I was at about [low / medium / high] capacity." 2. **Friction:** "One small thing that didn't feel great today was ____." (Optional — skip if there's nothing.) 3. **Ask:** "One thing that would help tomorrow is ____." (Optional.)

No debate, no solving, no rebuttal. Each person just transmits; the other just receives. Anything that needs real discussion gets flagged for the weekly closing (Chapter V), not litigated in the 90 seconds.

The weekly resync

Once a week, alongside closing the books:

- Skim the week's daily loop notes for anything that repeated more than twice.
- A repeated friction point is promoted to the RACI treaty (WP-03) or the tone filter (WP-09) rather than staying an ambient irritation.

Failure mode to watch for

If the daily loop starts being used to deliver criticism ("well, since we're doing check-ins...") instead of a plain status update, it will get avoided, and avoidance is what causes the drift this protocol exists to prevent. Keep the daily version strictly observational; save anything evaluative for the weekly resync.

Sign-off

	Committed to daily loop	Date
--	-------------------------	------

Partner A		
Partner B		

A NOTE ON THE CALCULATOR

CALC-01 — Solvency Read

The interactive Solvency Read calculator is a live, browser-based tool and isn't reproduced as a static page here. Access it any time at:

spreadloveandacceptance.com/workpapers/calculators/calc01-solvency.html